



Streaming ETL Patterns with DLT



Agenda

Streaming ETL Patterns with DLT

Lesson Name	Lesson Name
Lecture: Data Ingestion Patterns	Lecture: Data Modeling
Data Ingestion Patterns ADE 2.1 – Follow Along Demo – Auto Load to Bronze ADE 2.2 – Follow Along Demo – Stream from Multiplex Bronze	ADE 2.5 – Follow Along Demo – Data Modeling – SCD Type 2
Lecture: Data Quality Enforcement Patterns	Lecture: Streaming Joins and Statefulness
ADE 2.3 – Follow Along Demo – Data Quality Enforcement	ADE 2.6 – Follow Along Demo – Streaming Joins
ADE 2.4L – Streaming ETL Lab	



Data Ingestion Patterns



Why Do We Need These Patterns?

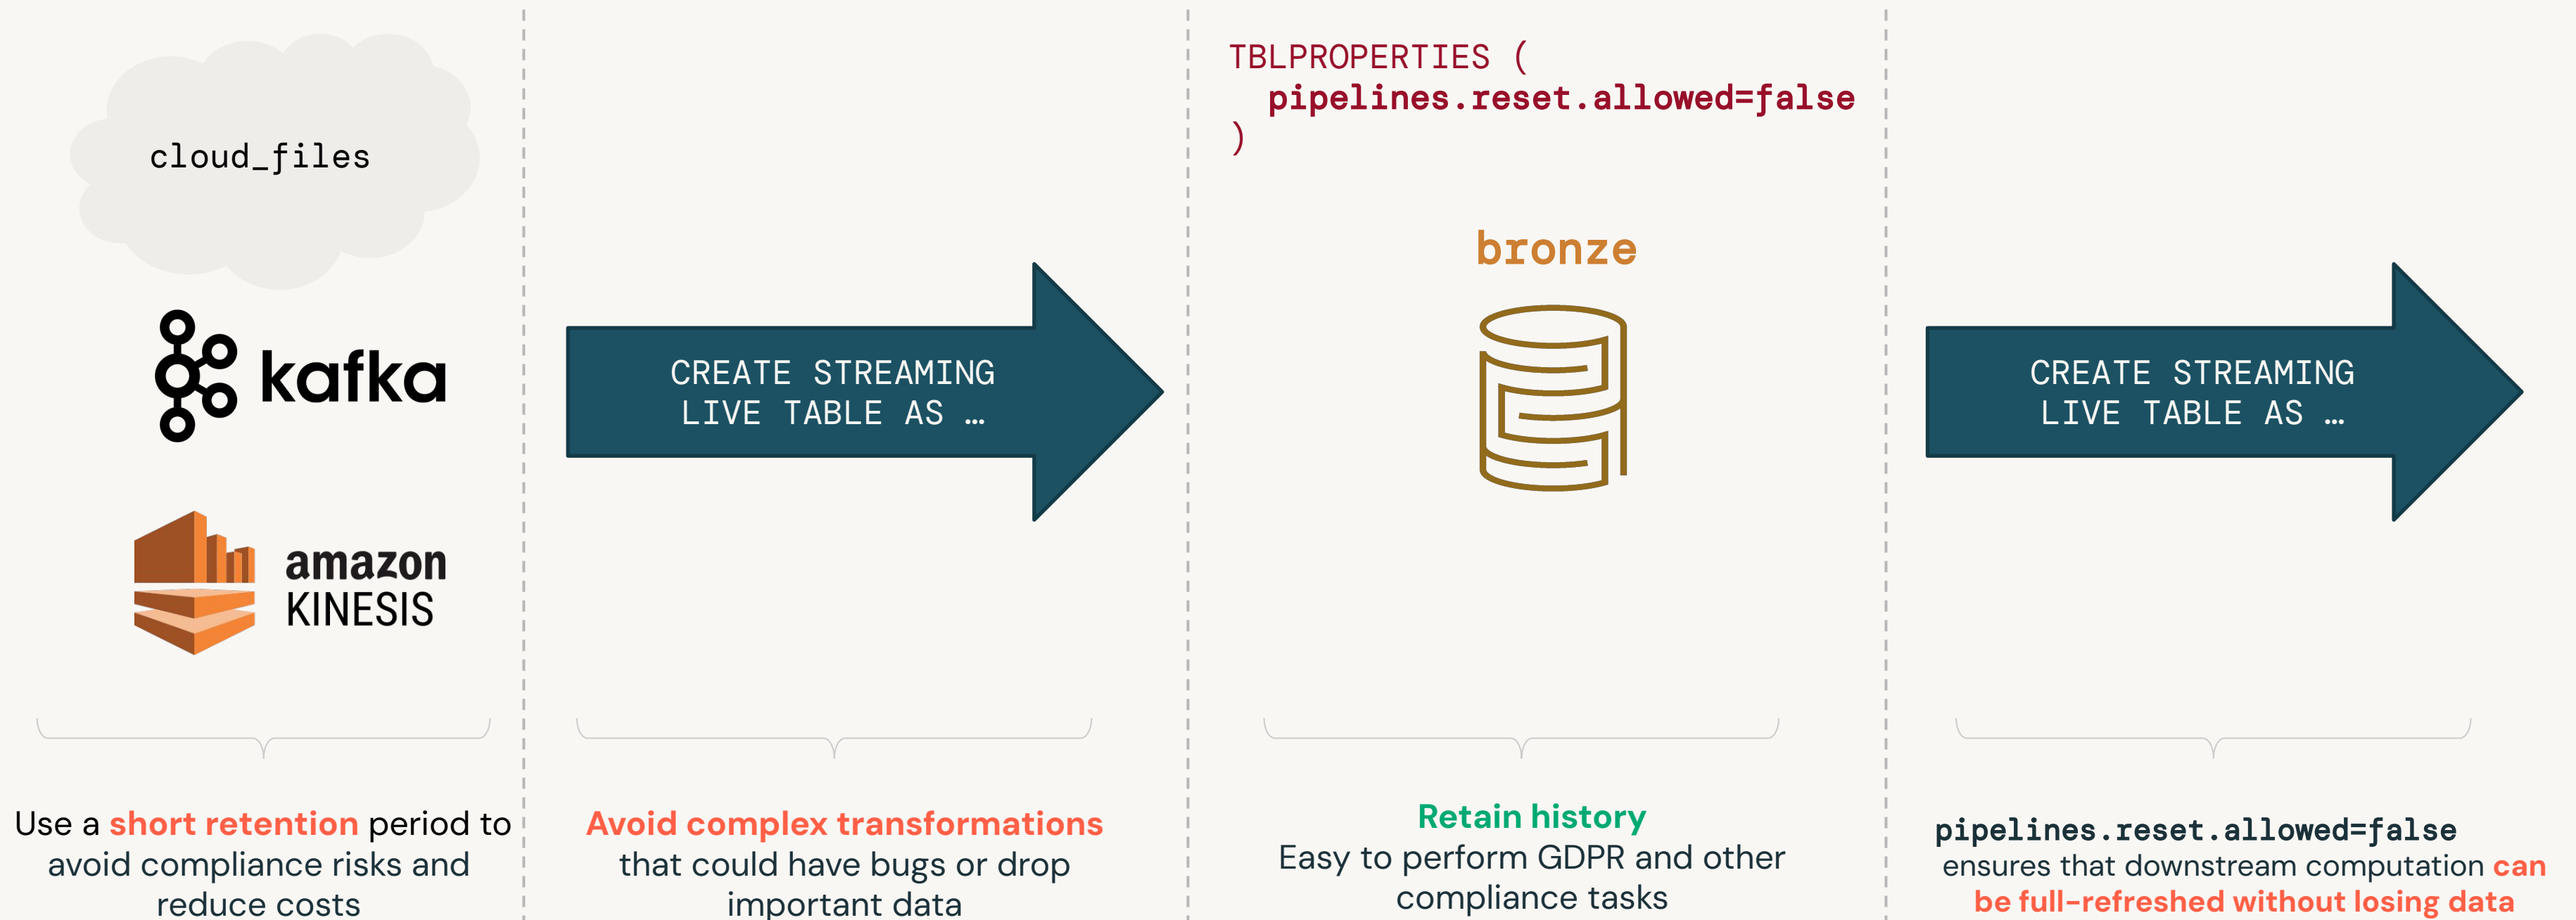
Limitations at Data Ingestion Stage

- Streaming sources like Kinesis, Kafka and EventHubs only retain data for a **limited** amount of time
- **Need for retention** – full history of data
 - Reprocessing raw data
 - Perform GDPR and compliance tasks
 - Recover data
- Need for a simple, **maintainable and scalable** architecture
- Keeping full history in the streaming source is **expensive**



Pattern 1: Use Delta for Infinite Retention

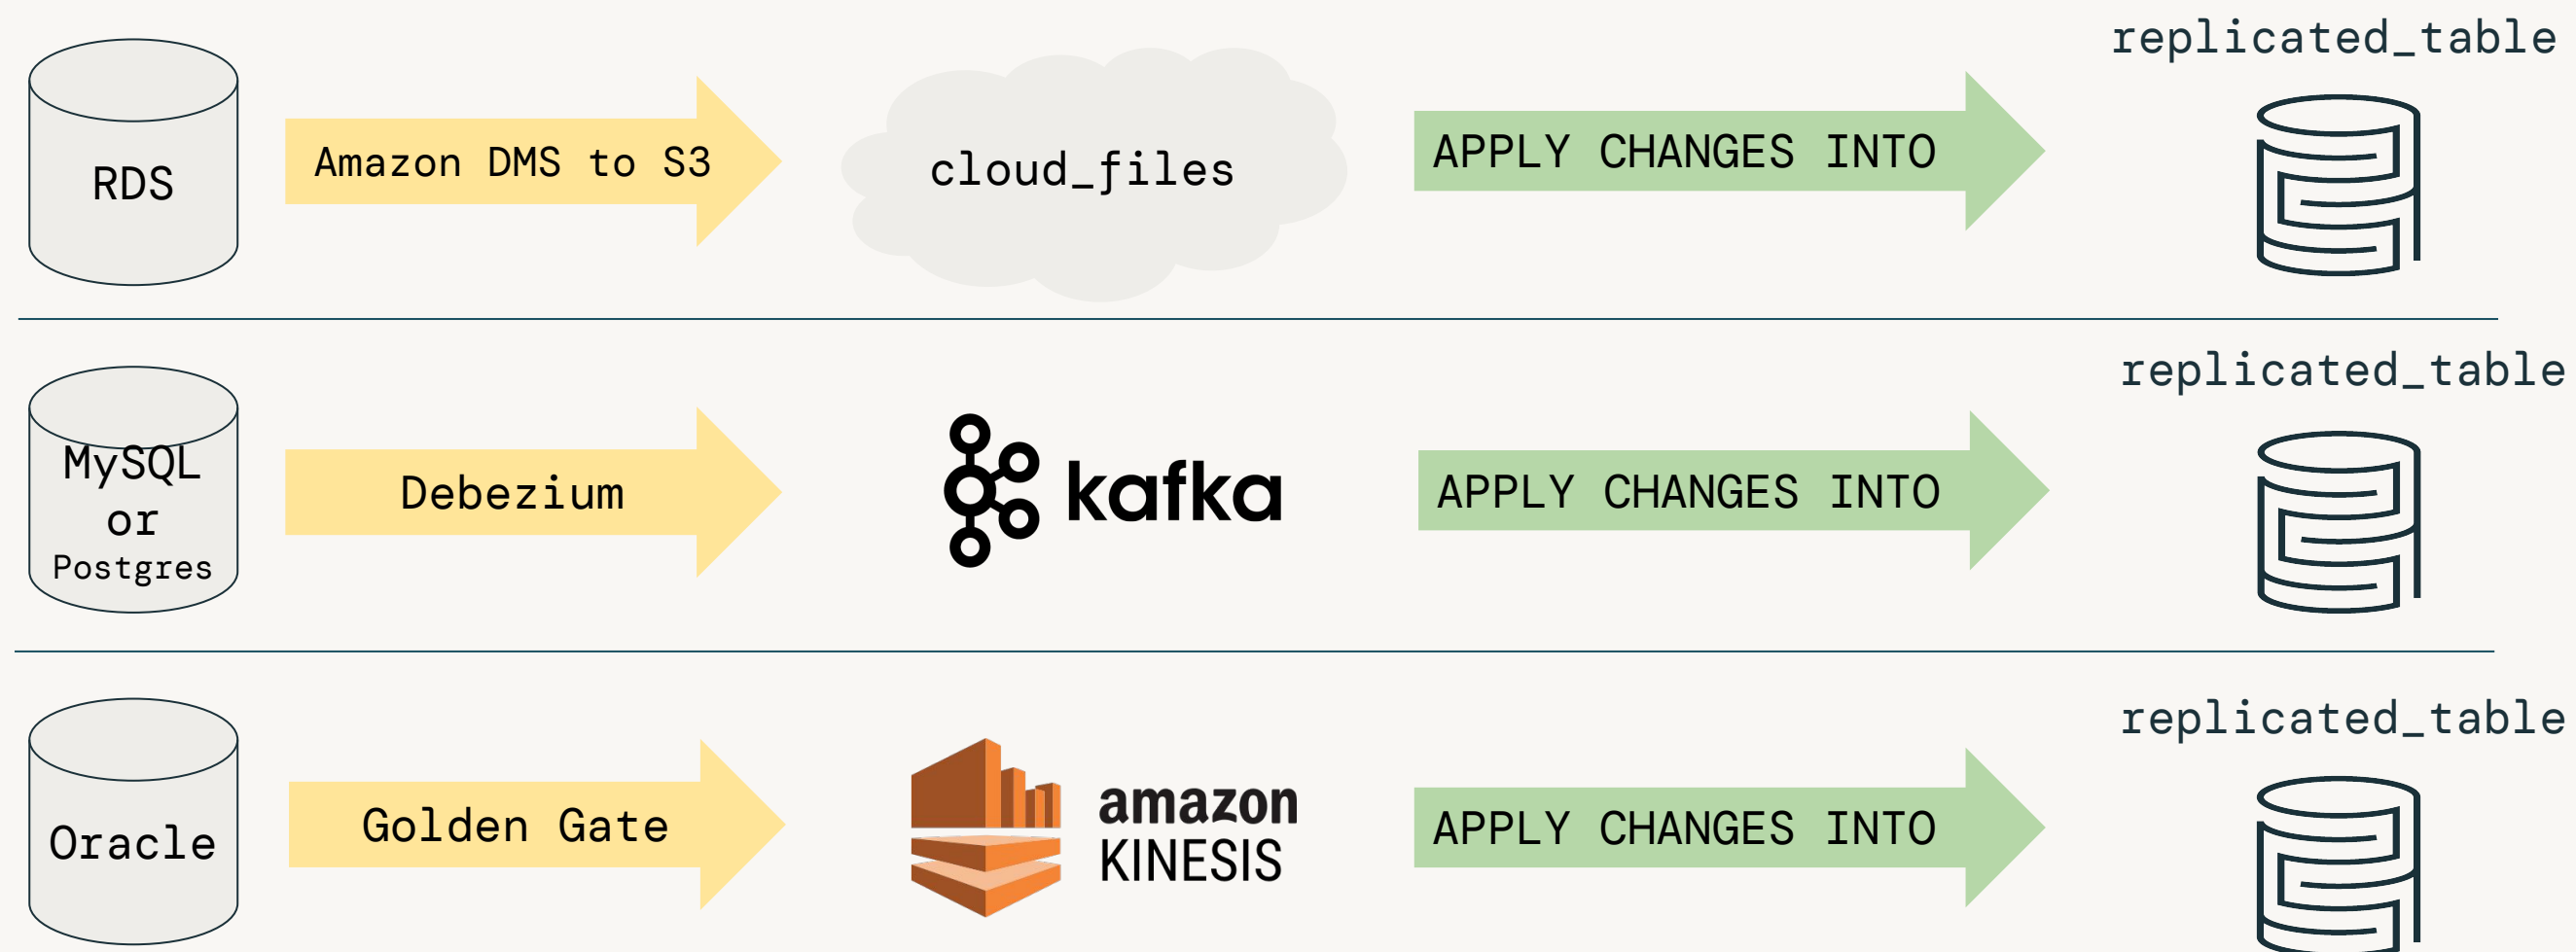
Delta provides cheap, elastic and governable storage for transient sources



Pattern 2: Up-to-date Replica with CDC

Maintain an up-to-date replica of a table stored elsewhere


- Use Change Data Capture (CDC) from RDMS and create replica as Delta
- A variety of 3rd party tools can provide a streaming change feed



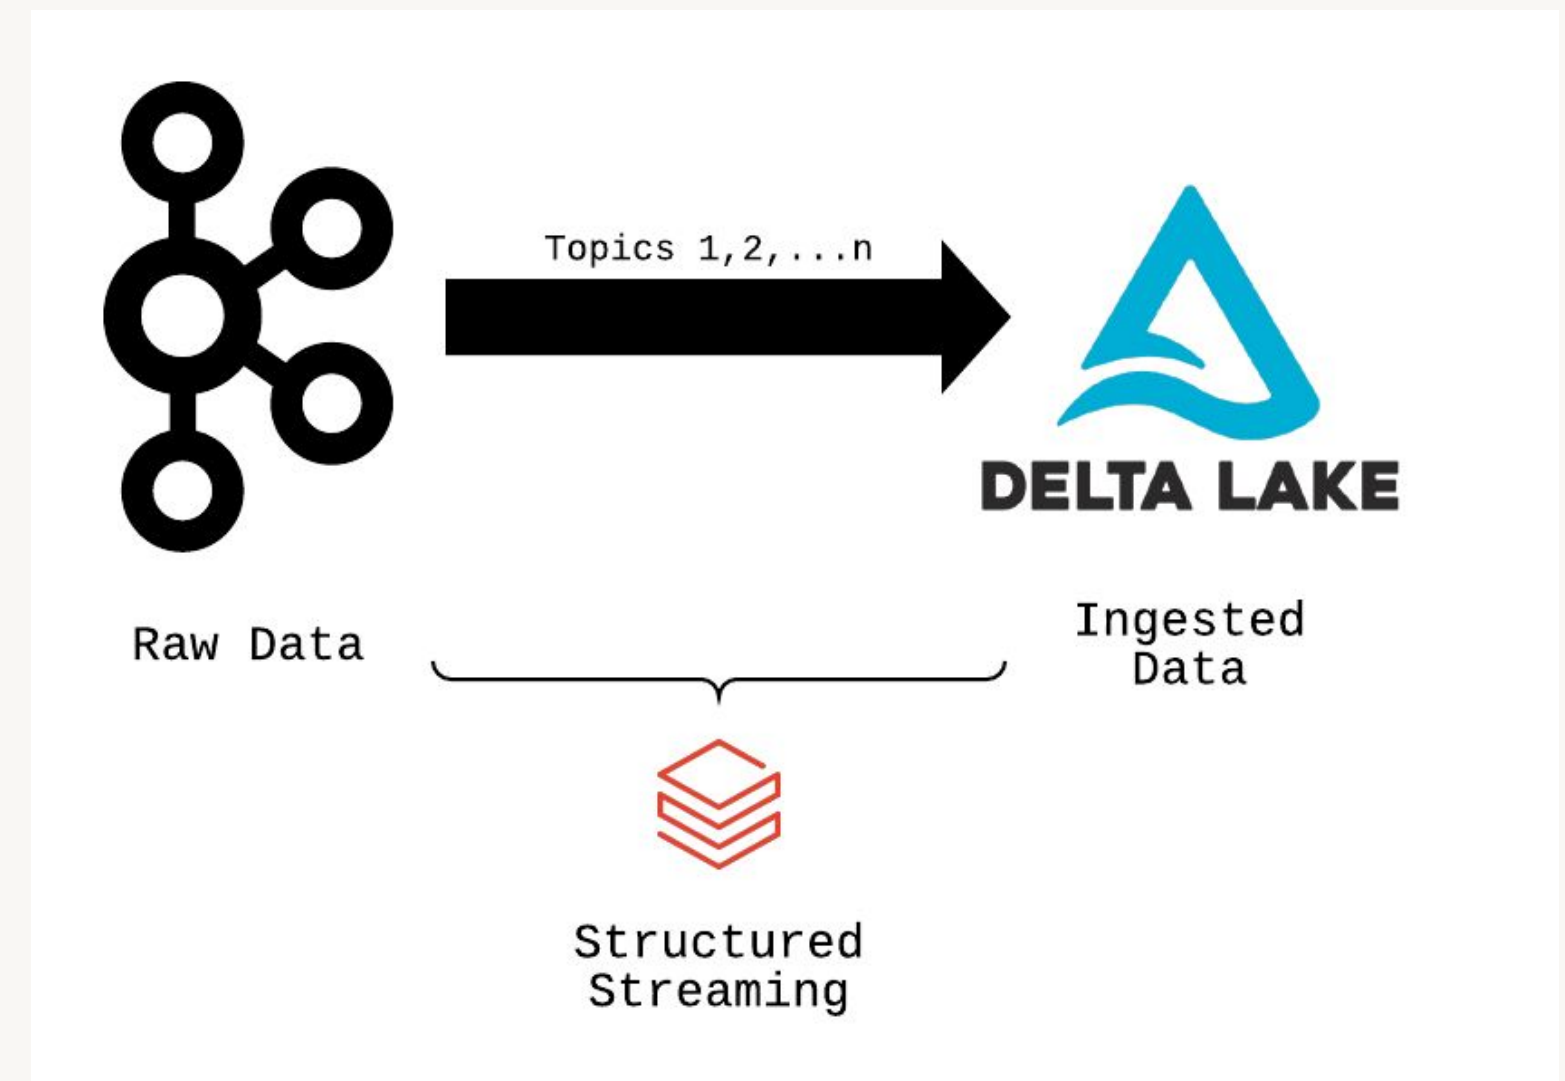
Pattern 3: Multiplex Ingestion

Multiplexing is used when a set of independent streams all share the same source

Simplex



Multiplex

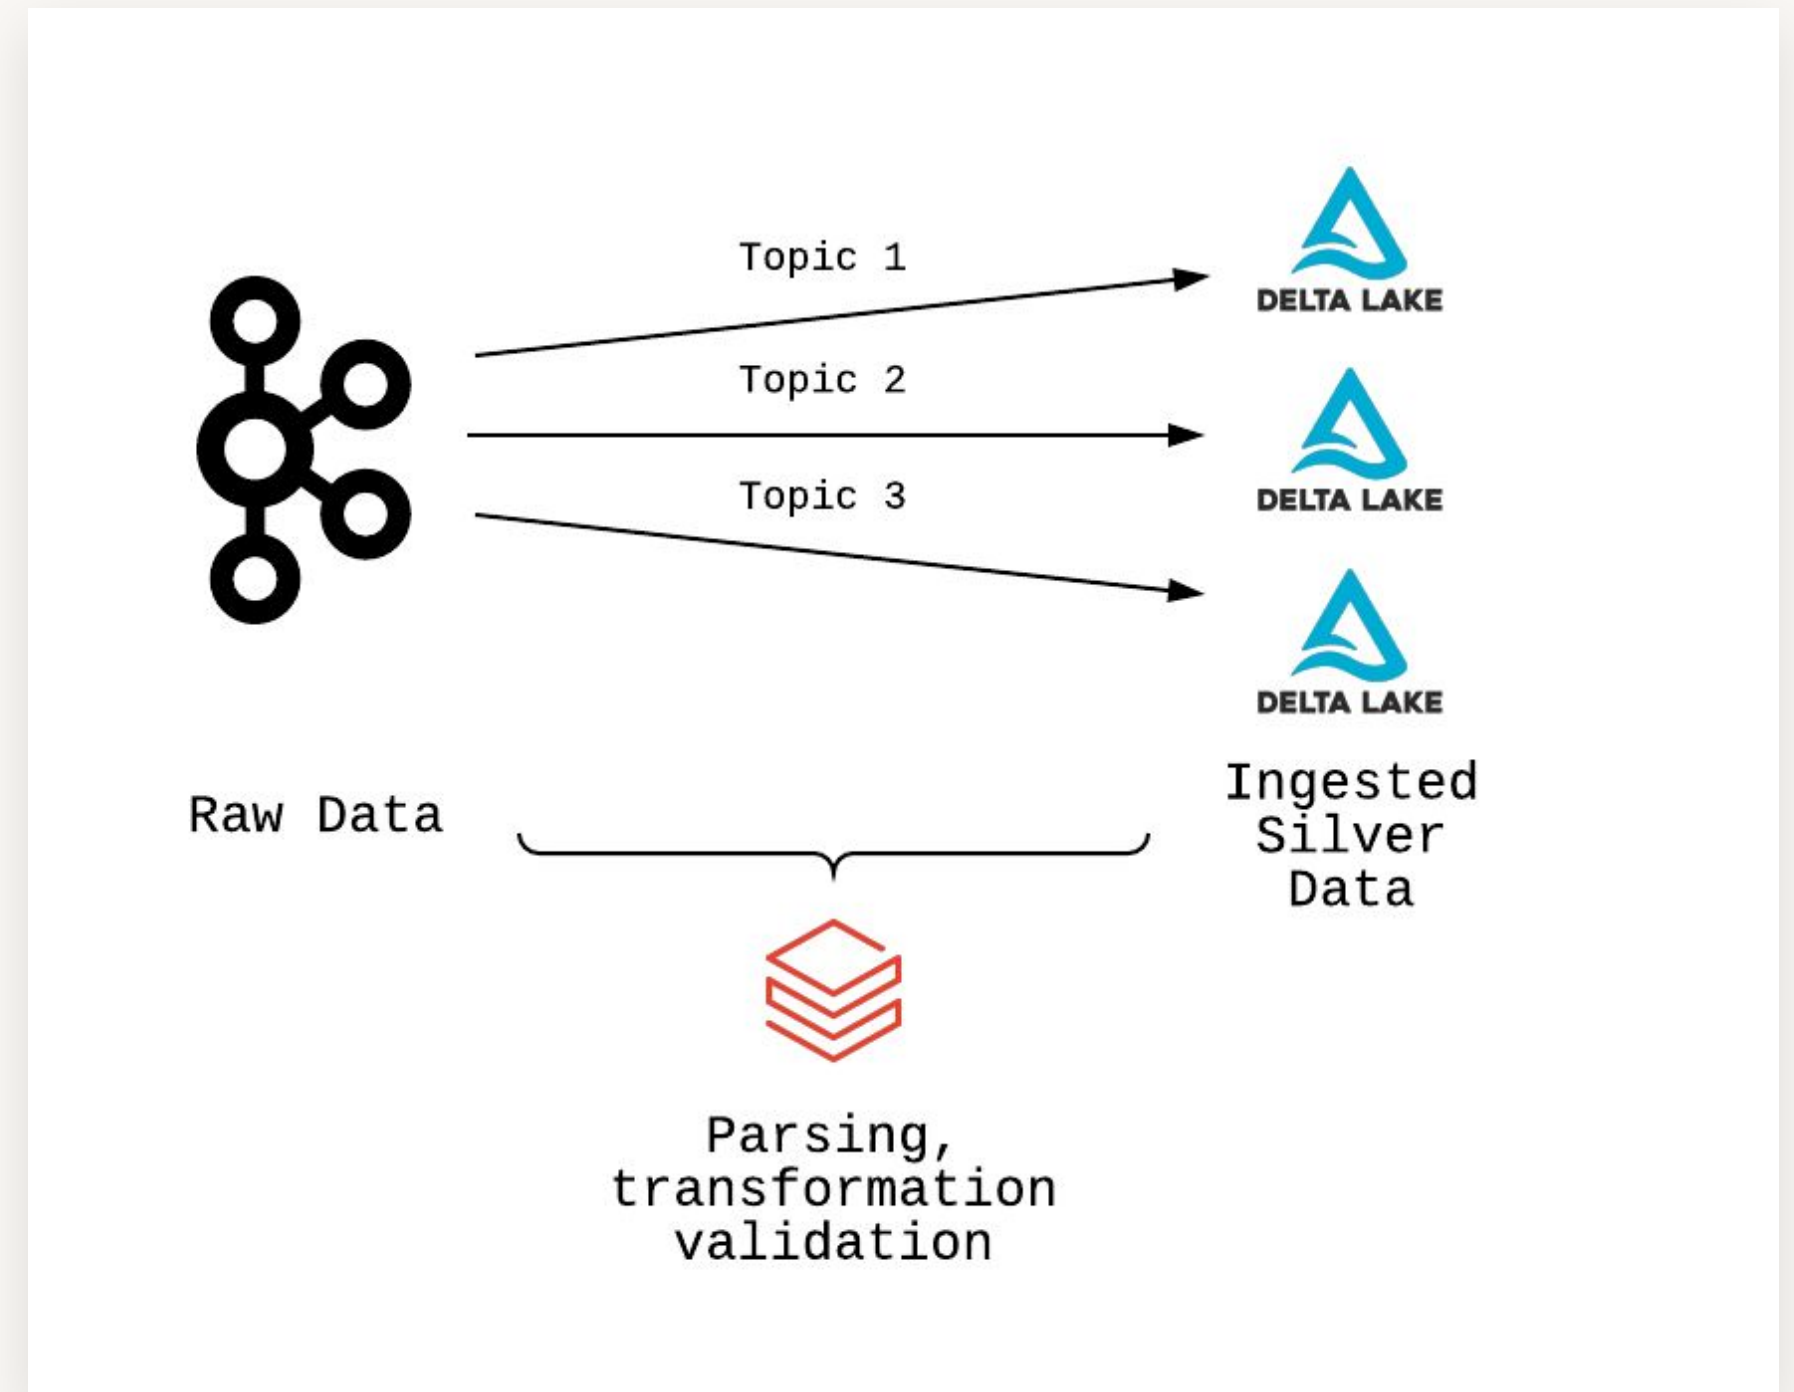


Pattern 3: Multiplex Ingestion

Anti-Pattern: Using Kafka as Bronze Table

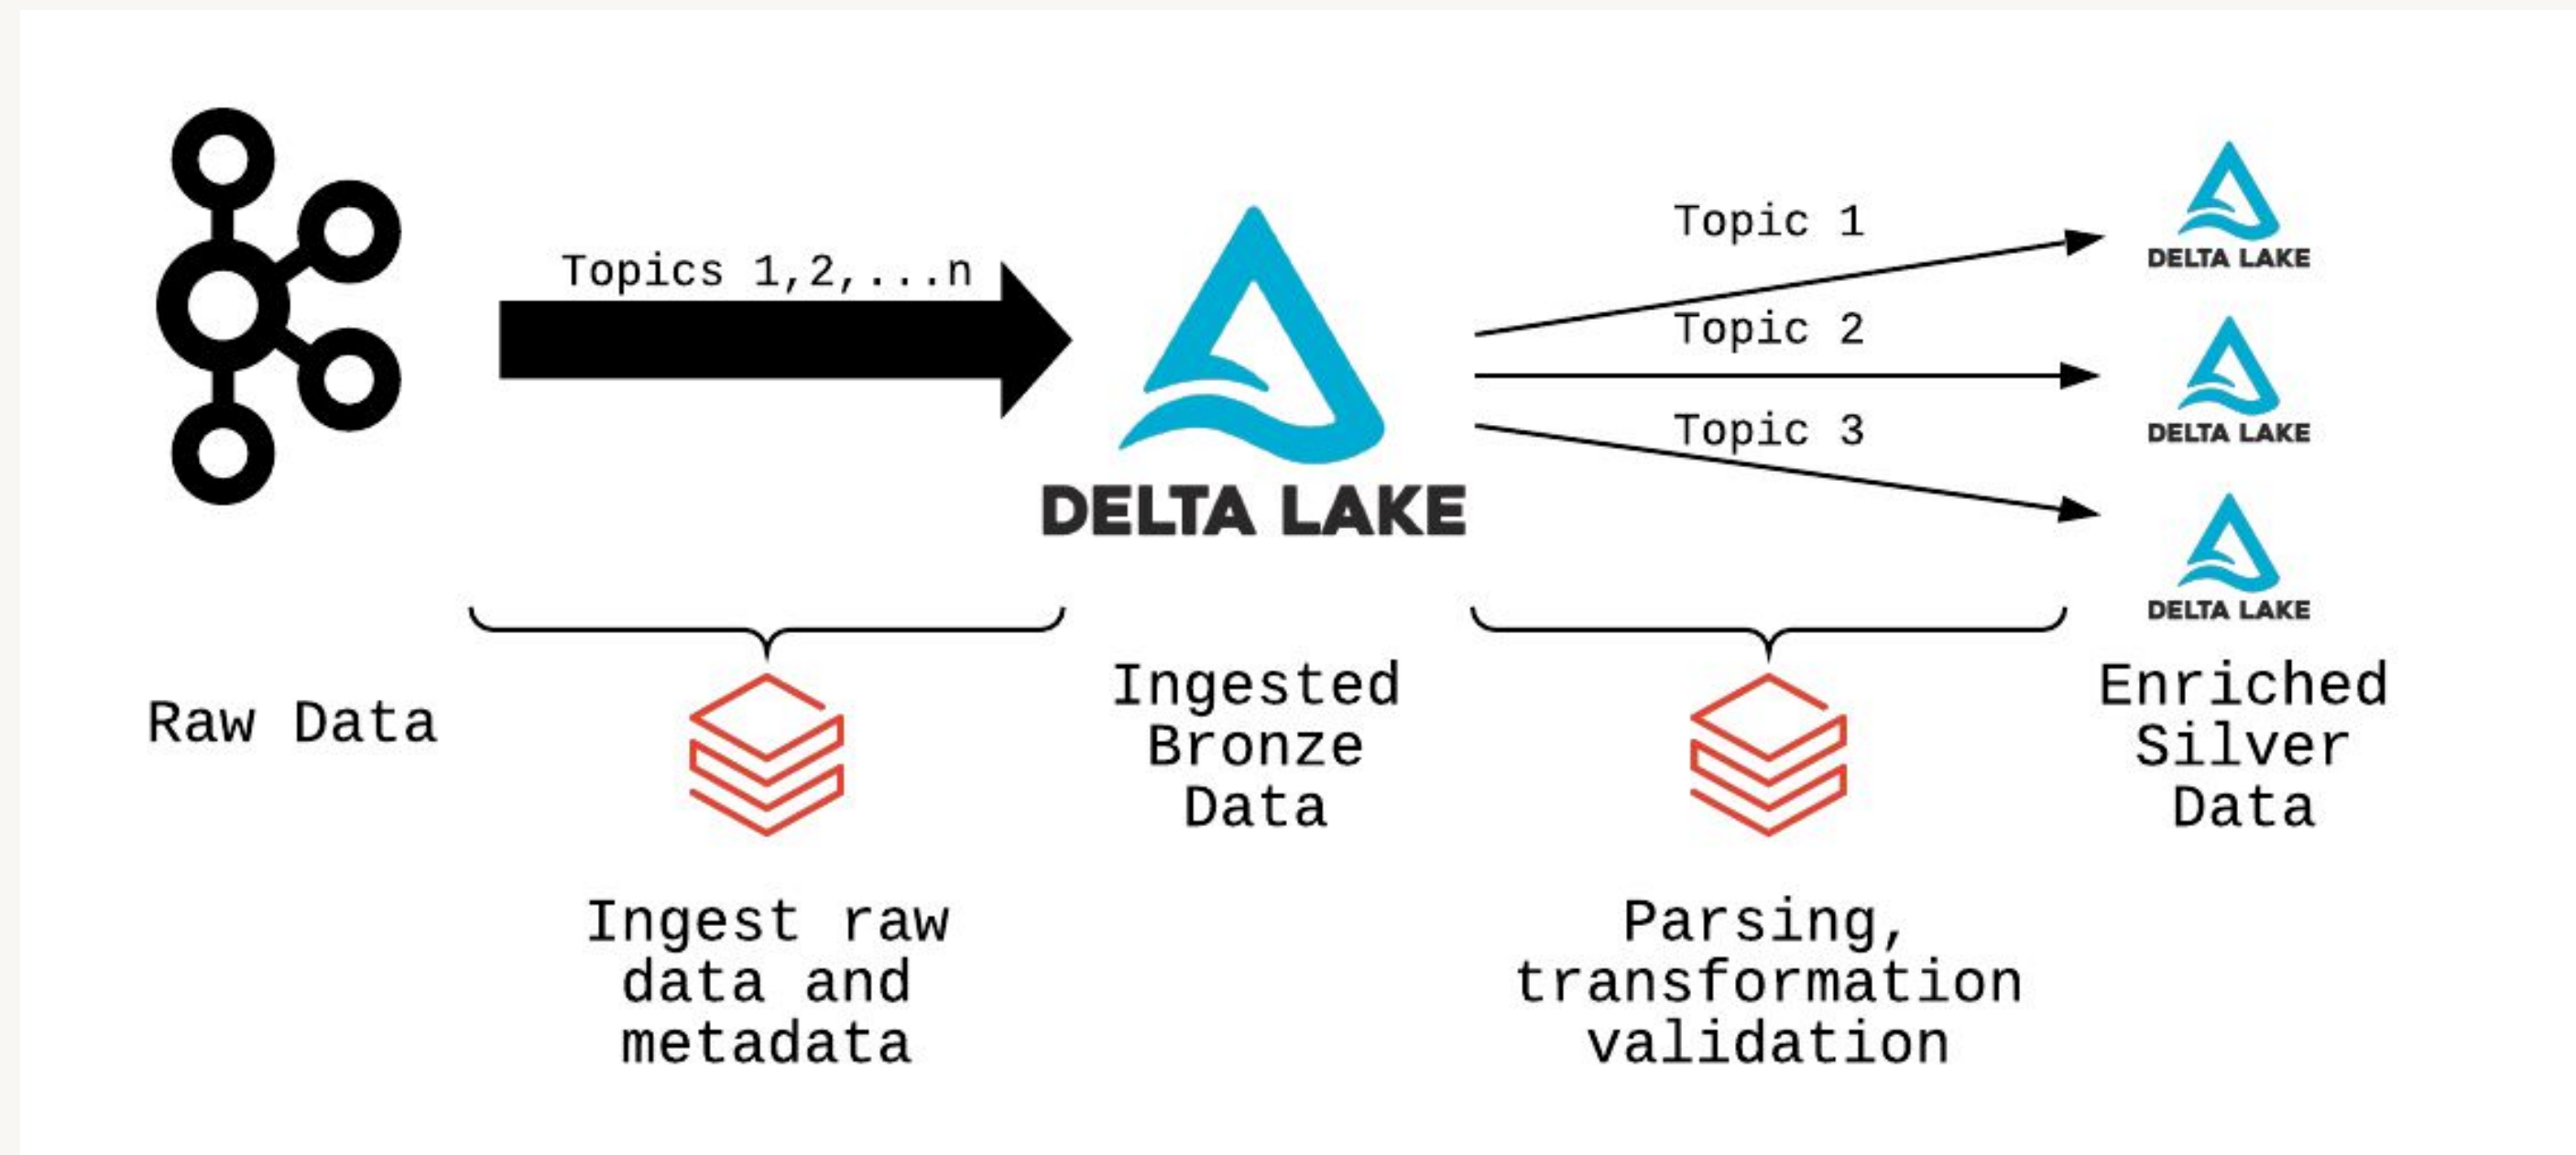
Don't use Kafka as Bronze Table:

- Data retention limited by Kafka; expensive to keep full history
- All processing happens on ingest
- If stream gets too far behind, data is lost
- Cannot recover data (no history to replay)



Pattern 3: Multiplex Ingestion Pattern

Multiplexing is used when a set of independent streams all share the same source



Demo: Auto Load to Bronze



Demo: Stream from Multiplex Bronze





Data Quality Enforcement Patterns



Silver Layer for Quality Enforcement

Silver Layer Objectives

- Validate data quality and schema
- Enrich and transform data
- Optimize data layout and storage for downstream queries
- Provide single source of truth for analytics



Schema Enforcement & Evolution

- Enforcement prevents bad records from entering table
 - Mismatch in type or field name
- Evolution allows new fields to be added
 - Useful when schema changes in production/new fields added to nested data
 - Cannot use evolution to remove fields
 - All previous records will show newly added field as Null
 - For previously written records, the underlying file isn't modified.
 - The additional field is simply defined in the metadata and dynamically read as null



Alternative Quality Check Approaches

- Add a “validation” field that captures any validation errors and a null value means validation passed.
- Quarantine data by filtering non-compliant data to alternate location
- Warn without failing by writing additional fields with constraint check results to Delta tables



Pattern: Quarantine Invalid Records

What if we want to save the records that violate data quality constraints for analysis?


Method 1: Create Inverse Expectation Rules



Limitations:

- Processes the data twice

Method 2: Add a validation status column and use for partitioning



Limitations:

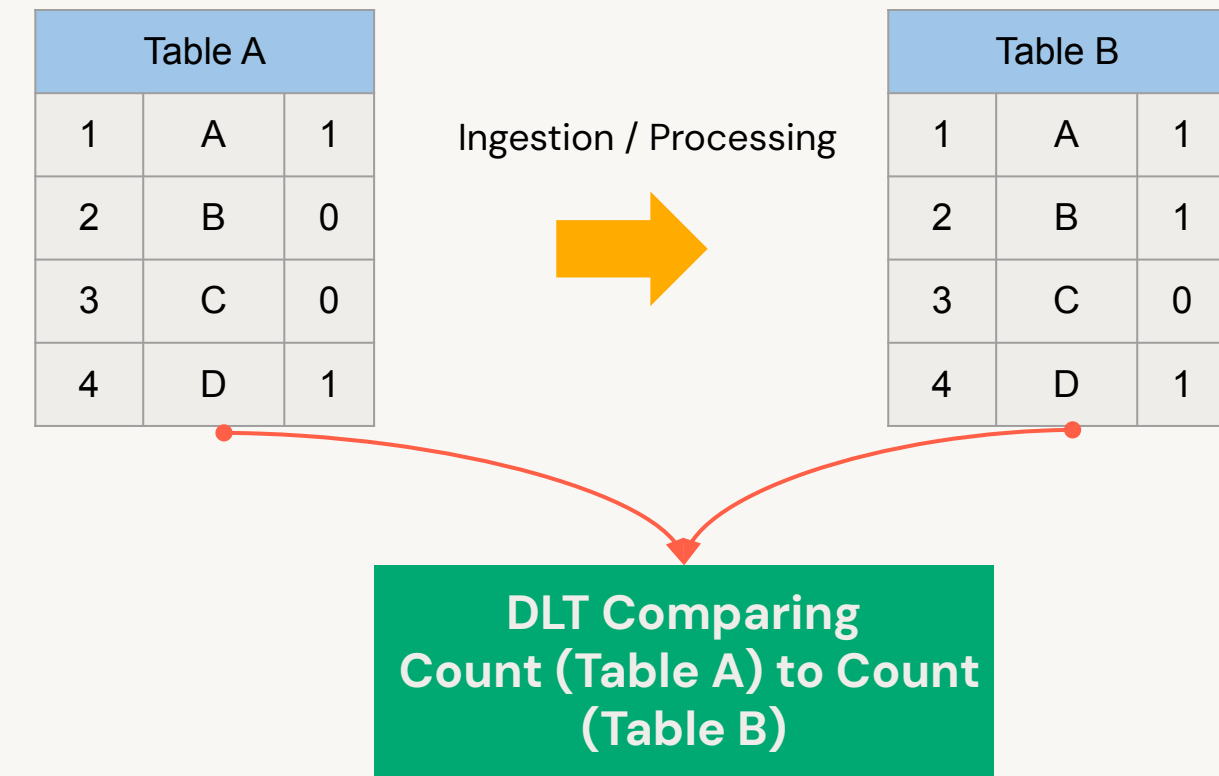
- Doesn't use expectations; data quality metrics are not available in the event logs or the pipelines UI.



Pattern: Verify Data with Row Comparison

Validate row counts across tables to verify that data was processed successfully without dropping rows.

- Solution:
 - Add an additional table to your pipeline that defines an expectation to perform the comparison.
 - The results of this expectation appear in the event log and the Delta Live Tables UI.



```
DLT.sql

CREATE OR REFRESH LIVE TABLE count_verification(
  CONSTRAINT no_rows_dropped EXPECT (a_count == b_count)
) AS SELECT * FROM
  (SELECT COUNT(*) AS a_count FROM LIVE.tbl_a),
  (SELECT COUNT(*) AS b_count FROM LIVE.tbl_b)
```



Pattern: Define Tables for Adv. Validation

Perform advanced data validation with DLT expectations

- Complex data quality checks examples;
 - A derived table contains all records from the source table
 - Guaranteeing the equality of a numeric column across tables
- Solution:
 - Define DLT using aggregate and join queries and use the results of those queries as part of your expectation checking.

```
-- Validates all expected records are present
in the "report" table

CREATE LIVE TABLE compare_tests(
  CONSTRAINT no_missing_records
  EXPECT (r.key IS NOT NULL)
)
AS SELECT * FROM LIVE.validation_copy v
LEFT OUTER JOIN LIVE.report r ON v.key = r.key
```



ADE 2.3 – Data Quality Enforcement



ADE 2.4L – Streaming ETL Lab






Data Modeling



Learning Objectives

By the end of this lesson, you should be able to:

-  1 Describe main concepts of dimensional modeling
-  2 Describe SCD tables and implementation with Delta Live Tables
-  3 Explain a common pipeline wherein a streaming data source joins to a static table.



Slowly Changing Dimensions in the Lakehouse




Dimensional Modeling

Fact Tables vs. Dimension Tables

- **Fact Tables:** Often contain a granular record of activities
- **Dimension Tables:** Often contain data may be updated or modified over time.

Modeling Guidelines:

- Denormalize dimension and fact tables
- Use conformed dimensions
- Use slowly changing dimensions as necessary



Dimensional Modeling

Fact Tables as Incremental Data

- Often is a time series
- No intermediate aggregations
- No overwrite/update/delete operations
- Often append-only operations



Slowly Changing Dimensions (SCD)

3 types of dimension tables

Type 0

- No changes allowed
- Tables are either static or append only
- Examples: static lookup tables, append-only fact tables

Type 1

- Overwrite but no history is maintained
- May contain recording of when record was entered, but not previous values
- Example: valid customer mailing address

Type 2

- Add a new row; mark old row as obsolete
- Strong history is maintained
- Example: tracking product price changes over time



Slowly Changing Dimensions (SCD)

3 types of dimension tables

Type 0 / Type 1

user_id	street	name
1	123 Oak Ave	Sam
2	99 Jump St	Abhi
3	1000 Rodeo Dr	Kasey

Type 2

user_id	street	name	valid_from	current
1	123 Oak Ave	Sam	2020-01-01	true
2	99 Jump St	Abhi	2020-01-01	false
3	1000 Rodeo Dr	Kasey	2020-01-01	false
2	430 River Rd	Abhi	2021-10-10	true
3	1000 Rodeo Dr	Casey	2021-10-10	true



SCD Type 2 with DLT

Keep a record of how values changed over time

```
APPLY CHANGES INTO LIVE.cities
FROM STREAM(LIVE.city_updates)
KEYS (id)
SEQUENCE BY ts
STORED AS SCD TYPE 2
```

__starts_at and __ends_at will
take on the type of the SEQUENCE
BY field (ts).

city_updates

```
{"id": 1, "ts": 1, "city": "Bekerly, CA"}
{"id": 1, "ts": 2, "city": "Berkeley, CA"}
```

cities

id	city	__starts_at	__ends_at
1	Bekerly, CA	1	2
1	Berkeley, CA	2	null



Applying SCD Principles to Facts

- Fact table usually append-only (Type 0)
- Can leverage event and processing times for append-only history

order_id	user_id	occurred_at	action	processed_time
123	1	2021-10-01 10:05:00	ORDER_CANCELLED	2021-10-01 10:05:30
123	1	2021-10-01 10:00:00	ORDER_PLACED	2021-10-01 10:06:30



ADE 2.5 – Data Modeling






Streaming Joins and Statefulness



The Components of a Stateful Stream

```
windowedDF =  
  (eventsDF  
    .groupBy(window("eventTime",  
                  "10 minutes",  
                  "5 minutes"))  
    .count()  
    .writeStream  
    .trigger(processingTime="5 minutes")  
  )
```



Statefulness vs. Query Progress

- Some operations are specifically **stateful** in that the results of processing earlier records from the stream affect the processing of later records.
 - Examples include deduplication, aggregation, and stream-stream joins
- Other transformations just need to store incremental query progress and are **not stateful**.
 - Examples include simple transformations and stream-static joins
- Progress and state are stored in **checkpoints** and managed by the driver during query processing.



Stream-Static Joins

Using **Dimension Tables** in Incremental Updates

- Delta Lake enables dynamic stream-static joins
- Each micro-batch captures the most recent state of the Delta table that is the static side of the join
 - This does not occur if the static side of the join is another format such as Parquet
- Allows modification of dimension while maintaining downstream composability

Note: Because Delta Lake does not enforce foreign key constraints, it is possible that joined data will go unmatched.



Streaming Queries are Not Stateful

Each input row is **processed only once**

A change to a streaming live table's definition **does not recompute** results by default:

```
CREATE STREAMING LIVE TABLE raw_data
AS SELECT a + 1 AS a a * 2 AS a
FROM cloud_files("/data", "json")
```

"/data"

{"a": 1}

{"a": 2}

{"a": 3}

{"a": 4}

raw_data

b
2
3
6
8



Streaming Joins are Not Stateful

Enrich data by joining with an **up-to date-snapshot** stored in Delta

A **change** to joined table snapshot **does not recompute** results by default:

```
CREATE STREAMING LIVE TABLE raw_data
AS SELECT *
FROM cloud_files("/data", "json") f
JOIN prod.cities c USING id
```

"/data"

{"a": 1} →

{"a": 1} →

raw_data

id	city
1	Bekerly, CA
1	Berkeley, CA

id	city
1	Bekerly, CA

Berkeley, CA

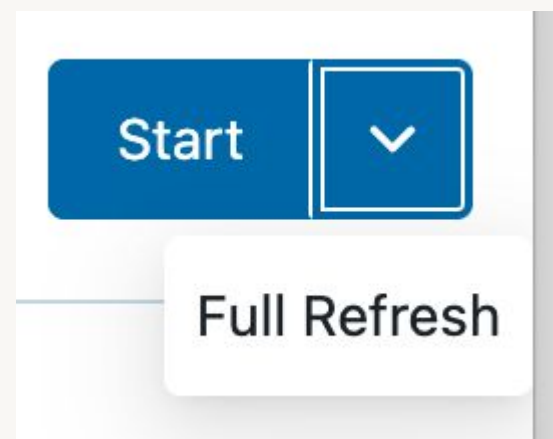


Clear State in DLT

Perform backfills after critical changes using **full refresh**

Full-refresh clears the table's data and the queries state, **reprocessing all the data**.

```
CREATE STREAMING LIVE TABLE raw_data
AS SELECT a * 2 AS a
FROM cloud_files("/data", "json")
```



"/data"

`{"a": 1}` →
`{"a": 2}` →
`{"a": 3}` →
`{"a": 4}` →

raw_data

b
2
3
6
8

After full-refresh

`{"a": 1}` →
`{"a": 2}` →
`{"a": 3}` →
`{"a": 4}` →

b
2
4
6
8



Stream-Static Join & Merge

- Join driven by streaming data
- Join triggers shuffle
- Join itself is stateless
- Control state information with predicate
- Goal is to broadcast static table to streaming data
- Broadcasting puts all data on each node

1. Main input stream

```
salesSDF = (  
    spark  
        .readStream  
        .format("delta")  
        .table("sales")  
)
```

2. Join item category lookup

```
itemSalesSDF = (  
    salesSDF  
        .join(  
            spark.table("items")  
                .filter("category='Food'"), #  
            Predicate  
                on=["item_id"]  
        )  
)
```



ADE 2.6 – Streaming Joins



Knowledge Check



Which of the following is considered a recommended best practice for ingesting streaming data?

Select one response.

- A. Use streaming live tables for raw data and streaming tables for bronze, silver, and gold quality data.
- B. Use streaming tables for bronze quality data and streaming live tables for silver and gold quality data.
- C. Use streaming live tables for bronze quality data and streaming tables for silver and gold quality data.
- D. Use streaming tables for raw data and streaming live tables for bronze, silver, and gold quality data.



A data engineer has data that needs to be updated. However, they need to have access to a recorded history of the information previously stored in the dataset before the update. Which of the following table types should the data engineer use for their data?

Select one response.

- A. Type 0
- B. Type 1
- C. Type 2
- D. Type 1 or Type 2



Which of the following operations can be performed on stateless tables to limit the state dimension?

Select one response.

- A. Stream-stream join
- B. Stream-static join
- C. Stateful aggregation
- D. Drop duplicates



Which of the following statements about fact tables and dimension tables are true?

Select two responses.

- A. Transactional guarantees and Delta Lake ensure that the newest version of a dimension table will be referenced each time a query is processed for incremental workloads.
- B. Joined data cannot go unmatched because of Delta Lake's foreign key constraint.
- C. Dimension tables contain a granular record of activities, while fact tables contain data that is updated or modified over time.
- D. Modern guidelines suggest denormalizing dimension and fact tables.



The following line of code is supposed to create a set of inverted rules for a quarantine table.

```
quarantine_rules = _____
```

Which of the following correctly fills in the blank?

Select one response.

- A. `{"invalid_record": f"NOT({ ' AND ' .join(rules.values()))" }`
- B. `{"invalid_record": f"&&({ ' ! ' .join(rules.values()))" }`
- C. `{"invalid_record": f"NOT({ ' OR ' .join(rules.values()))" }`
- D. `{"invalid_record": f"IF({ ' NULL ' .join(rules.values()))" }`

